

# Chapter 1

## Introduction



# Rendezvous Problem

Imagine a person trying to start two programs at the same time

- on two separate computers
- wanting them to communicate with each other



# Rendezvous Problem—Client/Server

The client/server model solves the rendezvous problem by determining that one side (the server) must start execution, and wait indefinitely for the other side (the client) to contact it

- Server
  - a program that waits for incoming communication requests from a client
  - When the server receives a communication from the client, the server provides some useful service to the client, and sends the client the results
- Client
  - a program that initiates communication with the server
  - The client has need of some service that the server can provide



# Client/Server

- Clients are often easier to build than servers
  - commonly require no special system privileges
- Servers often need to access data and/or routines or resources that are provided by an operating system and protected by the operating system
  - servers often need special system privileges



## Client/Server (cont'd)

- Servers are concerned with:
  - Authentication – verifying that the client is who it claims to be
  - Authorization – determining whether the given client is permitted to access any of the services the server supplies
  - Providing services
  - Data security – guaranteeing that a client is not allowed to access data that the client is not allowed to access, preventing data from being stolen



## Client/Server (cont'd)

Servers typically come in two different versions:

- Stateless – server does not save any information about the status of ongoing interaction with clients
- Stateful – server saves information about the status of ongoing interactions with clients



# What is Middleware?

According to Oracle (2016):

“Middleware is the software that connects software components or enterprise applications. Middleware is the software layer that lies between the operating system and the applications on each side of a distributed computer network. Typically, it supports complex, distributed business software applications.”



## What is Middleware? (cont'd)

This definition is from Techopedia (2016) :

“Middleware is a software layer situated between applications and operating systems. Middleware is typically used in distributed systems where it simplifies software development by doing the following:

- Hides the intricacies of distributed applications

- Hides the heterogeneity of hardware, operating systems and protocols

- Provides uniform and high-level interfaces used to make interoperable, reusable and portable applications

- Provides a set of common services that minimizes duplication of efforts and enhances collaboration between applications”





## What is Middleware? (cont'd)

This definition is from Apprenda (2016) :

“A simple middleware definition: software that connects computers and devices to other applications. It can also be referred to as the slash or connecting point in client/server.

Another way to define middleware is to say that it is software that acts as a liaison between applications and networks. The term is often used in the context of cloud computing, such as public or private cloud.”



## What is Middleware? (cont'd)

### Middleware

- is used to hide the lower level complexity of networks and operating systems from the application programmer
- can allow a client running on one kind of computer with one kind of operating system to talk to a server on a different kind of computer running a different operating system
- allows the definition of clearly defined interfaces to servers



# Sockets

- The Berkeley Software Distribution of UNIX back in 1983 defined an application interface called “sockets”
  - This eventually became a Portable Operating System Interface (POSIX) specification
  - So they’re commonly known as BSD sockets or as POSIX sockets



## Sockets (cont'd)

- There is a Windows version of sockets, called Winsock that was originally based on the POSIX sockets
  - However, there are quite a few differences in terms of include files, names of library routines, etc.
  - According to MSDN (2016) there are also some implementation differences based on differences in Windows compared to UNIX
- It is usually possible to port socket applications between Linux and Windows



## Sockets (cont'd)—Socket Data Structures

- With POSIX/BSD sockets, a socket is treated similarly to a file
  - Sockets are stored in the file descriptor table
  - thus, an application cannot have both a file descriptor and a socket descriptor with the same value



## Sockets (cont'd)—Socket Data Structures

- To perform communication, sockets choose a family of protocols to use
  - which family to use is defined in the *sockaddr* struct
    - which is defined in the file “sys/socket.h”



## Sockets (cont'd)—Socket Data Structures

```
struct sockaddr
{
    unsigned char  sa_len;
                    // length of address
    sa_family_t    sa_family;
                    // the address family
    char           sa_data[14];
                    // the address
};
```



## Sockets (cont'd)—Socket Data Structures

Several different address families are supported

- based on the kinds of protocols being used
- Some of the predefined address families included in the sockets specifications were for protocols that are not used a lot nowadays
  - remember that sockets have been around since the early 80s
- we will focus only on TCP/IP here





- The sockaddr\_in struct is defined in the file “netinet/in.h”
- The sockaddr\_in struct will contain the following members (at a minimum):
  - sa\_family\_t sin\_family;
  - in\_port\_t sin\_port;
  - struct in\_addr sin\_addr;



## Sockets (cont'd)—Socket Data Structures

- The `sin_family` is set to `AF_INET`
  - `AF_INET` is a constant value defined as the number 2
  - which selects TCP/IP using IPv4 addressing.
    - this corresponds to a protocol family called `PF_INET`
    - the internet protocol family for IPv6, which is also number 2
- `sin_port` is a port number
- `sin_addr` is a typical IP address



## Sockets (cont'd)—Socket Data Structures

- The `sockaddr_in6` struct is also defined in the file “`netinet/in.h`”.
- The `sockaddr_in6` struct will contain the following members (at a minimum):
  - `sa_family_t`     `sin6_family;`
  - `in_port_t`        `sin6_port;`
  - `uint32_t`          `sin6_flowinfo;`
  - `struct in6_addr` `sin6_addr;`
  - `uint32_t`          `sin6_scope_id;`



## Sockets (cont'd)—Socket Data Structures

- The `sin6_family` is set to `AF_INET6`
  - `AF_INET6` is a constant defined as the number 28 in base 10
  - selects TCP/IP using IPv6 addressing
  - corresponds to a protocol family called `PF_INET6`
    - the internet protocol family for IPv6, which is also number 28
- `sin6_port` value is used to store an IPv6 port number (same as IPv4 port numbers)
- `sin6_addr` stores the IPv6 address
- `sin6_flowinfo` is for IPv6 traffic class and flow information
- `sin6_scope_id` is the set of interfaces for a scope



## Sockets (cont'd)—Socket Data Structures

- `socklen_t`
  - a 32 bit integer
  - used to define the size of an address  
(among other things)



## Sockets (cont'd)—Socket Data Structures

- instead of using either `sockaddr_in` or `sockaddr_in6`, you could use *`sockaddr_storage`*
  - can be used to hold and pass around either IPv4 or IPv6 addresses
  - likely to have to do a lot of typecasting, since the various socket API calls (`connect`, `accept`, etc.) depend on the `sockaddr` struct



## Sockets (cont'd)—Socket Library Calls

- An application calls a “socket” to create a new socket for network communication
- returns a descriptor for the newly created socket
- arguments to the socket call include:
- protocol family
- type of service (stream or datagram).
  - stream means a data connection is being used where several blocks will be sent (this is how TCP/IP works)
  - datagram means a block of data is sent all by itself without a regular connection being established (this is how UDP works)



# Sockets (cont'd)—Socket Library Calls

## Summary of main socket calls:

- **Socket**—creates a descriptor for use in network communications
- **connect**—connect to a remote peer (client)
- **write**—send outgoing data across a connection
- **read**—acquire incoming data from a connection
- **close**—terminate communication and deallocate a descriptor
- **bind**—bind a local IP address and protocol port to a socket
- **listen**—set the socket listening on the given address and port for connections from the client and set the number of incoming connections from a client (backlog) that will be allowed in the listen queue at any one time
- **accept**—accept the next incoming connection (server)
- **recv**—receive the next incoming datagram
- **recvmsg**—receive the next incoming datagram (variation of **recv**)
- **recvfrom**—receive the next incoming datagram and record its source endpoint address
- **send**—send an outgoing datagram
- **sendmsg**—send an outgoing datagram (variation of **send**)
- **sendto**—send an outgoing datagram, usually to a prerecorded endpoint address
- **shutdown**—terminate a TCP connection in one or both directions
- **getpeername**—after a connection arrives, obtain the remote machine's endpoint address from a socket
- **getsockopt**—obtain the current options for a socket
- **setsockopt**—change the options for a socket





## Sockets (cont'd)—Socket Library Calls

- `write()`—writes to an operating system buffer, and blocks when the buffer is full
- When `read()` is used with the TCP/IP protocol, it extracts bytes and copies them to user's buffer. It blocks if there is no input data. When there is data, it fills the receiving buffer, then stops.
- When `read()` is used with the UDP protocol, it extracts one incoming UDP message. If the buffer cannot hold the entire message, `read()` fills the buffer and discards all remaining data in the UDP message
- Note that you can also use a `recv` instead of a `read`, and you could also do `recvmsg`, `recvfrom`, `sendmsg`, `sendto`. These would allow you to send a message without waiting for a response to tell whether the message was received by the other side or not
  - This is datagram type operation



## Sidebar Big Endian vs. Little Endian

- Let's choose the number 70A32C17 base 16, and assume it is stored starting at byte address 100.
- In big endian this would be stored as follows:
  - 100: 70
  - 101: A3
  - 102: 2C
  - 103: 17
- whereas in little endian this would be stored as follows:
  - 100: 17
  - 101: 2C
  - 102: A3
  - 103: 70



# Network Byte Order with Sockets

- Network byte order—a standard representation specifies a standard representation for binary integers in protocol headers in TCP/IP
  - specifies integers with the most significant byte first (big endian).
- Although the protocol software hides headers from application programs, a socket programmer has to understand network byte order
  - since some socket routines require arguments to be stored in network byte order
  - For example, the protocol port field of a `sockaddr_in` structure uses network byte order



# Network Byte Order with Sockets

- socket routines include several functions that convert integers between network byte order and the local host's byte order
  - programs must explicitly call the conversion routines
  - and always should call them, even when not necessary, for portability
- Short conversion routines – operate on 16 bit integers (unsigned):
  - Host to network short (htons)
  - Network to host short (ntohs)
- Long conversion routines – operate on 32 bit integers (unsigned):
  - Host to network long (htonl)
  - Network to host long (ntohl)



# Sockets (cont'd)—Flow of Operations

## Client Flow

```
socket(...);
```

```
connect (...);
```

## Server Flow

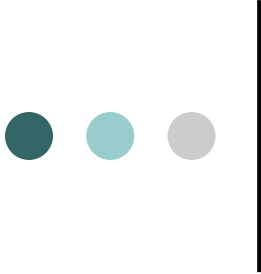
```
socket(...);
```

```
bind (IP address, protocol port);
```

```
listen(...);
```

```
while not end of time allowed  
{
```

```
    accept(...);
```



# Sockets (cont'd)—Flow of Operations (cont'd)

## Client Flow

write (...);

write(...)

close ();

## Server Flow (cont'd)

read(...);

read (...);

close();

}



```
server_socket_fd = socket(
    AF_INET,
    SOCK_STREAM,
    0
);
```



## Sockets (cont'd)—Simple Socket Server (cont'd)

```
bind( server_socket_fd,  
      (struct sockaddr *) &server_addr,  
      sizeof(server_addr))
```

```
listen(server_socket_fd, 3);
```





## Sockets (cont'd)—Simple Socket Server (cont'd)

```
client_conn_fd = accept( server_socket_fd,  
                        (struct sockaddr *) &client_addr,  
                        (socklen_t *) &client_addr_size  
                        );
```



## Sockets (cont'd)—Simple Socket Server (cont'd)

```
client_message_length = read( client_conn_fd,  
                             message_from_client,  
                             255  
                             );
```



## Sockets (cont'd)—Simple Socket Client

```
inet_pton(AF_INET,  
IP_address,  
&serv_addr.sin_addr  
)
```



## Sockets (cont'd)—Simple Socket Client (cont'd)

```
connect(    socket_fd,  
           (struct sockaddr *) &serv_addr,  
           sizeof(serv_addr)  
        )
```



# What are Remote Procedure Calls?

- The first attempts at middleware were things like sockets and Open Network Computing (ONC) Remote Procedure Calls (RPCs)
- With Remote Procedure Calls (such as ONC RPC), the idea is that you treat a call across the network the same as you would treat a call to a local procedure



# What are Remote Procedure Calls?

- RPCs are functionally-oriented because they came along before the object-oriented paradigm was widely used in industry
  - Functionally-oriented means they focused on procedure calls—data was separate
- Note, however, that how various technologies that employ Distributed Objects work is often also called, in a generic way, a remote procedure call



# What are Remote Procedure Calls?

- RPCs are functionally-oriented because they came along before the object-oriented paradigm was widely used in industry
  - Functionally-oriented means they focused on procedure calls—data was separate
- Note, however, that how various technologies that employ Distributed Objects work is often also called, in a generic way, a remote procedure call



# What are Remote Procedure Calls?

ONC RPCs are a well-known RPC-based technology

- originally from Sun Microsystems
- date from the mid-1980s
- the remote procedures that are to be called across the network are defined in a file with a “.x” extension
- Then a program called “rpcgen” is used to translate this .x file into handlers in the C language
  - a stub file is created for the client side and later linked with client code
  - a skeleton file is created for the server side and later linked with server code





# Synchronous vs. Asynchronous

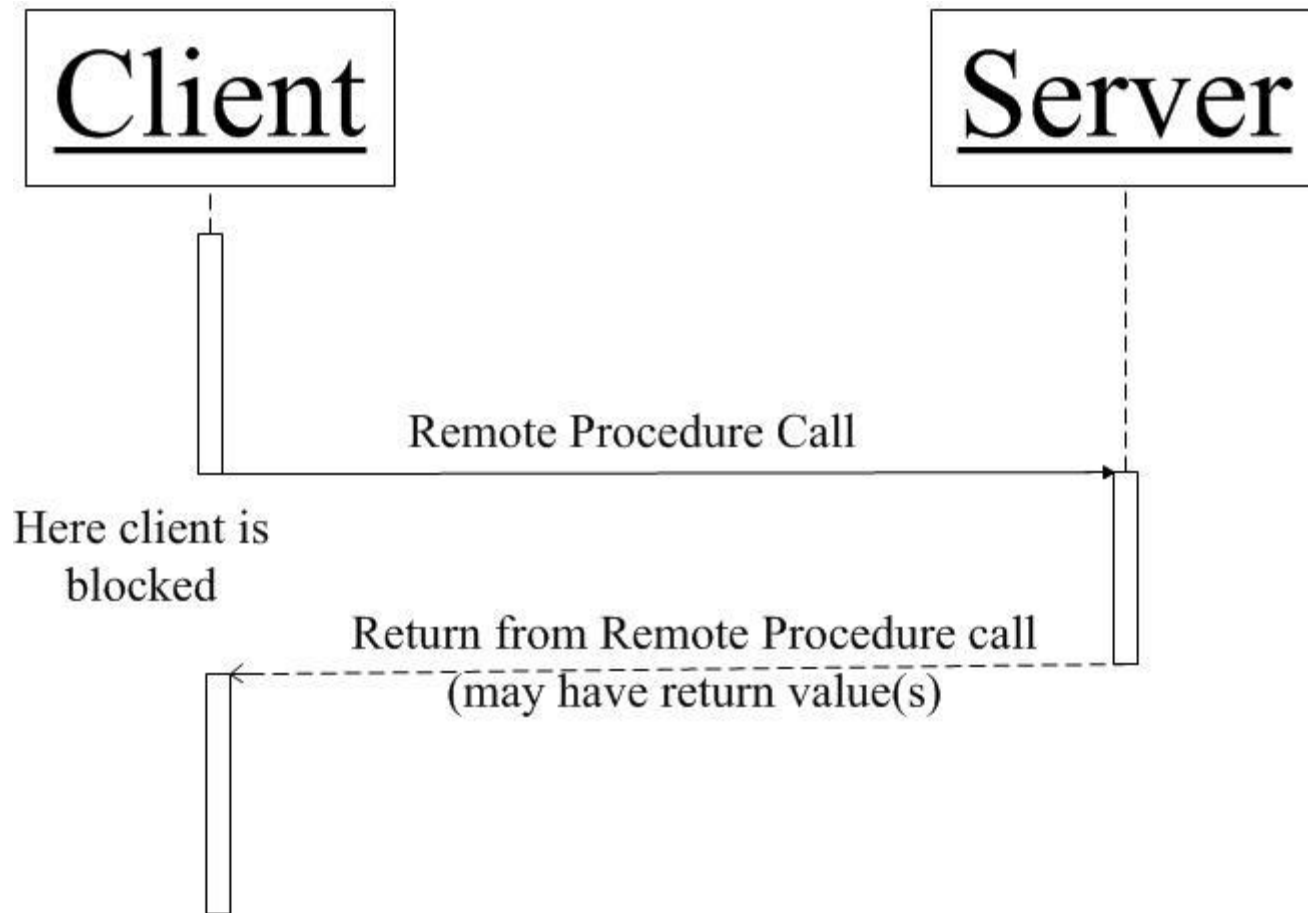
In synchronous communication

- the client calls the server
- then the client blocks and waits for the server to finish

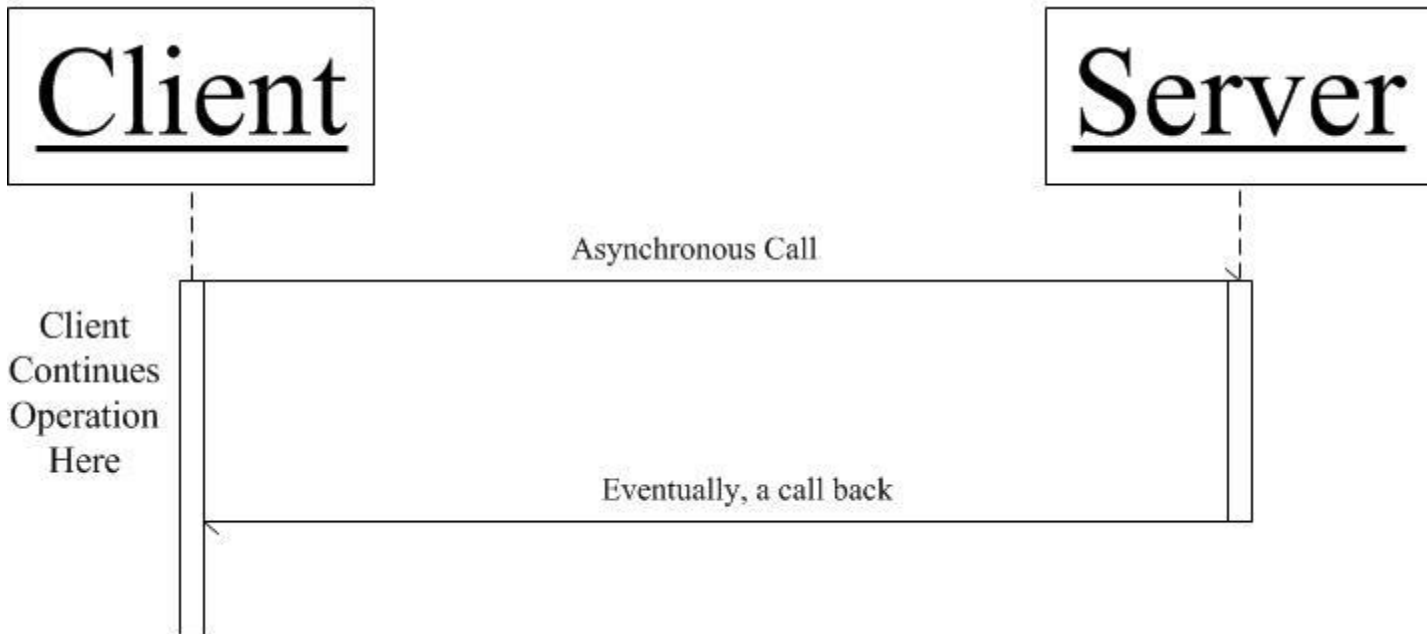
In asynchronous communication

- the client calls the server
- instead of blocking and waiting for the server's response, it goes on about its business

# Synchronous vs. Asynchronous (cont'd)— Synchronous



# Synchronous vs. Asynchronous (cont'd)— Asynchronous—Callback Version





# What are Distributed Object-Oriented Components?

In the object-oriented paradigm, of course data and procedure calls are encapsulated inside an object

So distributed objects just refer to objects that are located on different computers



# What are Distributed Object-Oriented Components?

Generally speaking, component-based software engineering consists of combining loosely coupled components into systems



# What are Distributed Object-Oriented Components?

There have been arguments about whether there is, in fact, a difference between component-based systems and object-oriented systems

- objects in an object-oriented system focus on modeling real world situations
- components in a component-based system are focused solely on combining existing components into systems
- but where do the existing components come from in the first place goes the other argument



# What are Distributed Object-Oriented Components?

Gomaa (2011) defines a distributed component as follows:

“a distributed component is a concurrent object with a well-defined interface, which is a logical unit of distribution and deployment. A well-designed component is capable of being reused in applications other than the one for which it was originally developed.”



# What is Message Oriented Middleware?

Message oriented middleware (MOM)

- purpose is to provide a middleman in between the client and the server
  - may do additional processing, perhaps Quality of Service-type processing
- the client talks to the middleman
- the server talks to the middleman,
- the client doesn't have to talk directly to the server
- the server doesn't have to talk directly to the client.
- nearly always uses a message queue to store messages between message producer and message consumer
- asynchronous






# What is Message Oriented Middleware?

Message oriented middleware (MOM)

- purpose is to provide a middleman in between the client and the server
  - may do additional processing, perhaps Quality of Service-type processing
- the client talks to the middleman
- the server talks to the middleman,
- the client doesn't have to talk directly to the server
- the server doesn't have to talk directly to the client.
- nearly always uses a message queue to store messages between message producer and message consumer
- asynchronous

# What is Message Oriented Middleware? (cont'd)






# What are Service Oriented Architectures?

According to the Open Group a *Service Oriented Architecture* (SOA) is an architectural style that supports service orientation

- Service orientation is a way of thinking in terms of the outcomes of services, and how they can be developed and combined
- a *service* is a repeatable business activity that can be logically represented, the Open Group gives the example: “check customer credit”
- a service is self-contained
- a service may be composed of other services
- consumers of the service treat the service as a black box

# What are Service Oriented Architectures? (cont'd)





## What are Service Oriented Architectures? (cont'd)

- Services may be located on different computers
- Each service has a well defined, well documented interface
- Services are independent and may be reused by different applications
- Services are loosely coupled with any of the calling applications
- It is common to implement service oriented architectures with web services
  - It is possible to alternately use other implementations, such as distributed object components



# What are Web Services?

- Web services are applications that typically expect to make use of the world wide web to provide application services
  - We distinguish the “world wide web” from the “internet” The internet is connected using the TCP/IP protocols
  - The world wide web is a collection of information in the form of web pages, that are connected using hypertext (clickable) links



# What are Web Services?

- Web services are applications that typically expect to make use of the world wide web to provide application services
  - We distinguish the “world wide web” from the “internet” The internet is connected using the TCP/IP protocols
  - The world wide web is a collection of information in the form of web pages, that are connected using hypertext (clickable) links



## What are Web Services? (cont'd)

- There are two kinds (architectural styles) of technologies that are commonly used for web services:
  - non-RESTful web services implemented with WSDL and SOAP technologies that treat HTTP as an independent lower protocol layer
  - the RESTful web services that use HTTP directly





# What is Cloud Computing?

- IBM (2016) defines cloud computing as follows:  
“Cloud computing, often referred to as simply “the cloud,” is the delivery of on-demand computing resources—everything from applications to data centers—over the Internet on a pay-for-use basis.”
- Applications running on the cloud often employ a service oriented architecture



# What is Cloud Computing?

- Public cloud:
  - instead of doing your own computing and storing your data on the computer on your desktop yourself
    - you hire a company to do the computing and store the data on their big computer servers that you access via the web
- Private cloud:
  - instead of having computing and data storage on employees' desks
  - the company can buy its own big servers
    - then the employees do their computing and store their data on those servers
- Hybrid cloud:
  - Part of a company's computing is done in house
  - Part of a company's computing is done by a public cloud



# What is Cloud Computing?

- There are three different paradigms for cloud computing:
  - Infrastructure as a Service (IaaS)
    - you or your company pays a cloud provider for computing resources
    - You provide your own operating system and application software
  - Platform as a Service (PaaS)
    - you or your company pays a cloud provider for an environment provided by the company that provides everything you need in order to develop and run your applications
      - This environment includes operating system, development tools, web site hosting, among other things
  - Software as a Service (SaaS)
    - you or your company pays a cloud provider for the use of their software application



# Environmental Monitoring Project

- A program used to illustrate various technologies in this textbooks
- Sensors will monitor Lake Guntersville for
  - Temperature
  - air quality
  - water quality
- Data is sent back to the University of Alabama in Huntsville (UAH) for analysis



# Sailboat Marina Project

- A program used to illustrate various technologies in this textbooks
- A web-based application will be used to manage a large sailboat marina that
  - employs a database to keep track of:
    - Sailboats
    - Owners
    - Whether owners have paid slip rental or not
  - Handles sailboat maintenance requests from owners