Lecture 8

Neutron, Data, Restful Interface, Restful Python Interface

OpenStack Internals The primary sub-parts of the OpenStack Cloud software:

	Service	Comments
nova	Compute	Manages pools of computer resources
glance	Image	Disk and server images
keystone	Identity	Common authentication system
horizon	Dashboard	GUI
cinder	Block Storage	Persistent block storage devices
swift	Object Storage	Object and file management
neutron	Networking	IP addresses, VLANS, etc.
trove	Database	Provisions relational and non-relational DBs
heat	Orchestration	Launches composite cloud applications based on templates, specify relationships between resources
ironic	Bare metal provisioning	Provisions and turns on and off bare metal machines instead of virtual machines
ceilometer	telemetry	Billing system
sahara	Elastic Map Reduce	Handles Hadoop clusters
zaqar	Multiple Tenant Cloud Messaging	Can be used in SaaS to send messages between multiple components of the SaaS

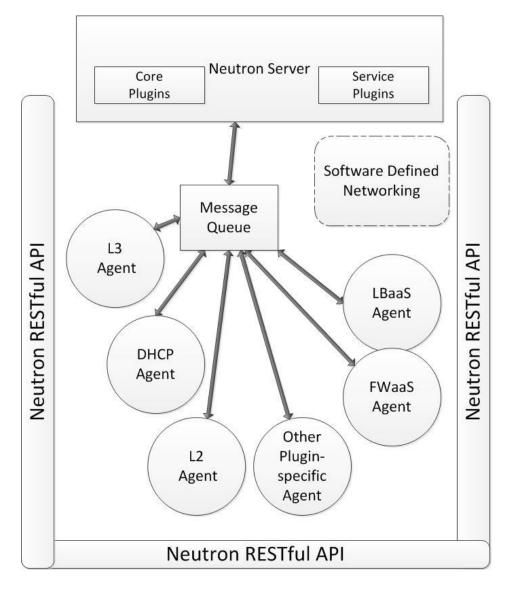
OpenStack Neutron

- Neutron provides considerable functionality that was not provided by nova-network.
- Neutron allows users to configure their own network topology including:
 - multi-tier networks (web tier, application tier),
 - their own network (how to assign IP addresses, etc.)
- It also allows use of advanced network services, including services intended to improve security and quality of service.

OpenStack Neutron (cont'd)

- When Neutron is being used (and not novanetwork), there is a portion of Neutron that runs on the controller node as a python daemon
 - It provides network APIs and passes messages to Neutron plugins

o openStack Neutron (cont'd)--Overview



OpenStack Neutron (cont'd)

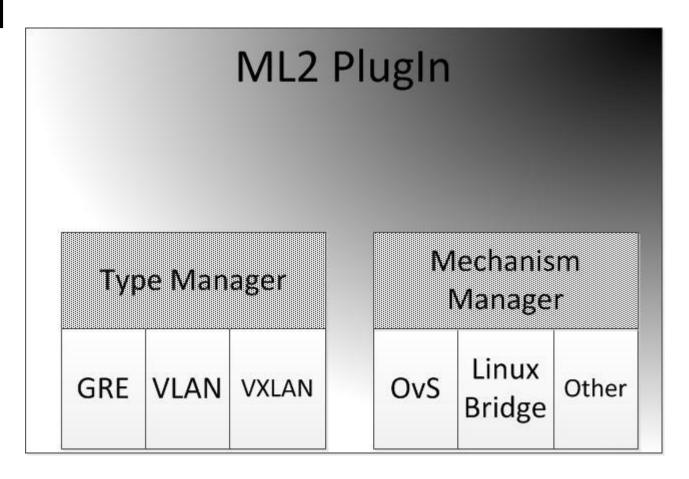
- OpenStack Neutron components communicate with each other using AMQP and a message queue.
- Neutron uses various plugin agents to implement different protocols or functions.
 There are two kinds of plugins:
 - Core plugins
 - Service plugins

OpenStack Neutron (cont'd)

- Core plugins implement basic network management including
 - IP address management (DHCP agent) and
 - L2 networking (ML2 plugin),
- Service plugins implement additional services including (among others):
 - load balancing (Load Balancing as a Service or LBaaS)
 - VPNs (VPNaaS)
 - firewalls (FWaaS)
 - metering

- OpenStack Neutron (cont'd)—ML2 Plugin
 - The Modular Layer 2 (ML2) plugin is a core Neutron plugin that allows Neutron to work with different kinds of layer 2 networks.
 - ML2 has types and mechanisms:
 - type is the type of network being used
 - the mechanism is the technology being used to implement the network

OpenStack Neutron (cont'd)—ML2 Plugin (cont'd)



- OpenStack Neutron (cont'd)—ML2 Plugin
 - In the previous figure, we see three kinds of network types listed (other kinds can be added), these are:
 - GRE (Generic Routing Encapsulation)
 - VLAN (Virtual Local Area Network)
 - VXLAN (Virtual Extensible Local Area Network)

- OpenStack Neutron (cont'd)—ML2 Plugin
 - Generic Routing Encapsulation (GRE) is a tunneling protocol used across an IP network to encapsulate many other protocols.
 - Tunneling protocols are often used to allow a different (foreign) protocol to run over a network that does not support that particular protocol
 - They are also often used to provide services that the underlying network does not support

OpenStack Neutron (cont'd)—ML2 Plugin

 A Virtual Local Area Network allows partitioning of a network such that users (in this case, VMs) on one VLAN cannot see the traffic on a different VLAN

- OpenStack Neutron (cont'd)—ML2 Plugin
 - In legard to ML2 mechanisms, OvS stands for Open vSwitch:
 - Open vSwitch is a well known open source method for creating virtual switches to connect to VMs in a software defined network
 - Open vSwitch employs Openflow, which is a software defined network standard
 - Linux bridges are an alternative method to create a virtual switch to connect to a VM

- OpenStack Neutron (cont'd)—ML2 Plugin
 - The purpose of a Virtual Extensible Local Area Network is to improve scalability of networking in the cloud by extending the address space available to create VLANs.
 - To do this, it overlays a layer 2 network on top of a layer 3 network.
 - Like a VLAN, a VM on a VXLAN can only see other VMs on that VXLAN, they can't see the traffic on other VXLANs.

- OpenStack Neutron (cont'd)—ML2 Plugin
 - Note that Neutron allows connecting to full blown SDN controllers, such as Open Daylight or Floodlight
 - SDN controllers can be used to provide a view of the complete network topology and to monitor the network state, both physical and virtual.
 - SDN controllers can also manage changes to the SDN.
 - This can be done either as a separate service plugin, or through the ML2 plugin (or both).

OpenStack Internals The primary sub-parts of the OpenStack Cloud software:

	Service	Comments
nova	Compute	Manages pools of computer resources
glance	Image	Disk and server images
keystone	Identity	Common authentication system
horizon	Dashboard	GUI
cinder	Block Storage	Persistent block storage devices
swift	Object Storage	Object and file management
neutron	Networking	IP addresses, VLANS, etc.
trove	Database	Provisions relational and non-relational DBs
heat	Orchestration	Launches composite cloud applications based on templates, specify relationships between resources
ironic	Bare metal provisioning	Provisions and turns on and off bare metal machines instead of virtual machines
ceilometer	telemetry	Billing system
sahara	Elastic Map Reduce	Handles Hadoop clusters
zaqar	Multiple Tenant Cloud Messaging	Can be used in SaaS to send messages between multiple components of the SaaS

OpenStack Data Storage

- There are three kinds of storage with OpenStack:
 - Ephemeral
 - Block Storage
 - Object Storage

- OpenStack Data Storage (cont'd)Each running Virtual Machine receives some ephemeral storage,
 - it's used to store the operating system of the image, and local data
 - The lifetime of the data in this storage is the lifetime of the Virtual Machine:
 - that is, the ephemeral data is no longer preserved when the Virtual Machine is terminated.

OpenStack Data Storage (cont'd)

Block storage is performed by Cinder,

Object Storage is performed using Swift.

- OpenStack Data Storage (cont'd)--Cinder
 - Cirder was originally part of Nova (called novavolume).
 - Cinder attaches storage volumes to a virtual machine
 - A storage volume is a storage area with its own file system
 - The server that is running Cinder can have storage volumes located on physical disks attached to the Cinder server itself, or can be located on other physical disks.

OpenStack Data Storage (cont'd)—Cinder (cont'd)

- o Cirder has three basic services:
 - cinder-scheduler
 - schedules the appropriate volume service
 - cinder-volume
 - manages block storage devices—does load balancing between volumes
 - cinder-backup
 - allows backup of volumes

- OpenStack Data Storage (cont'd)—Swift
 - OpenStack Swift stores data as binary objects (can have associated metadata) that are retrieved and written using HTTP commands (GET, PUT, etc.)
 - originally created by Rackspace, and part of the original OpenStack Austin release
 - Swift stores objects on object servers.
 - Swift is best used for storing unstructured data, such as email, images, audio, and video, etc.

OpenStack Data Storage (cont'd)—Swift (cont'd)

- Swift organizes object data into a hierarchy.
- Your account is the top level of the hierarchy, and you own the information associated with that account (project or tenant).
- Inside the account (project or tenant) can be containers (namespace for objects).
- Within one account, the names of containers must be unique.
- The containers may have associated Access Control Lists (ACLs) that specify who can have access to the objects within the container (individual objects do not have associated ACLs).

OpenStack Data Storage (cont'd)—Swift (cont'd)

- Swift has the concept of a "ring".
- A ring keeps track of where data is stored in the cluster
 - it is used to map names to the physical locations of objects (physical device on physical node)
- There are rings for:
 - accounts (projects or tenants)
 - containers
 - each separate storage policy

OpenStack Data Storage (cont'd)—Swift (cont'd)

- For data backup, rings are divided into zones, within which data is replicated.
- By default, three replicas of data objects are created and stored in a separate zone (to protect against hardware failures).
- A zone can consist of a single disk drive or a server in either the local data center or a different data center

- There are several different ways to access the OpenStack RESTful interface. Four of the most popular are:
- 1) the OpenStack Dashboard..
- 2) using cURL to send appropriate HTTP commands
- 3) using python with the python libraries provided with OpenStack
- 4) use the command line clients.

- Different OpenStack RESTful interfaces listen at different port numbers on the OpenStack url
 - See Table 13.10 text for some important OpenStack port numbers
- There are two versions of the Keystone interface commonly used today, version 2 and version 3.
 - Version 3 keystone url: http://localhost:5000/v3
 - Version 2 keystone url: http://localhost:5000/v2.0

- Different OpenStack RESTful interfaces listen at different port numbers on the OpenStack url.
- There are two versions of the Keystone interface commonly used today, version 2 and version 3.
 - Version 3 keystone url: http://localhost:5000/v3
 - Version 2 keystone url: http://localhost:5000/v2.0

- In the following slides, we will see one example of how to authenticate using each one of the different ways to access the OpenStack RESTful interface
- For how to start instances, etc., see the textbook and its associated code

φpenStack RESTful Interface—cURL

Keystone authentication using cURL command:

curl -i -H "Content-Type: application/json" \

-d @my_credentials \

http://localhost:5000/v3/auth/tokens

OpenStack RESTful Interface Interface— Credentials file in JSON format

```
"auth": {
     "identity": {
          "methods": [
                "password"
            "password": {
                 "user": {
                "name": "admin",
                 "domain": {
                     "id": "default"
                "password": "mypassword"
```

OpenStack RESTful Interface Interface— Keystore Authentication in python

```
from keystoneauth1.identity import v3
from keystoneauth1 import session
from keystoneclient.v3 import client
auth = v3.Password(
      user domain name='default',
      username='admin',
      password='mypassword',
      project domain name='default',
      project_name='admin',
      auth_url='http://localhost:5000/v3'
mysess=session.Session(auth=auth)
keystone=client.Client(session=mysess)
return keystone
```

OpenStack RESTful Interface Interface— Common Client

- Using the Keystone version 3 API
 - To authenticate, set the environment variables as follows:

```
export OS_AUTH_URL=http://localhost:5000/v3
export OS_USERNAME=admin
export OS_PASSWORD=mypassword
export OS_USER_DOMAIN_NAME=default
export OS_IDENTITY_API_VERSION="3"
export OS_PROJECT_NAME="admin"
```

Then type: openstack image list

OpenStack RESTful Interface Interface— Separate Keystone Client

Assuming you're using a devstack install:

In the devstack directory is a file called openrc.

Type:

source openrc

This will set various environment variables to default values so that you can access the openstack stuff. Alternately, set various variables, as follows:

OpenStack RESTful Interface Interface— Separate Keystone Client

- Since you set yourself up as the admin user, you can do some of the following commands:
 - To view all tenants (projects):
 - keystone tenant-list
 - To display all current users:
 - keystone user-list

OpenStack RESTful Interface Interface— Separate Keystone Client

- Since you set yourself up as the admin user, you can do some of the following commands:
 - To view all tenants (projects):
 - keystone tenant-list
 - To display all current users:
 - keystone user-list

OpenStack Python Interface

To enter the python environment, type: python

To exist the python environment, type; Ctrl-D

Authentication plugins are a way to extend how OpenStack performs authentication

OpenStack clients access these plugins in order to authenticate

Various standard "identity plugins" are provided with Keystone, this includes:

V2 identity plugins

V3 identity plugins

V2 identity plugins

- defined in module keystoneclient.auth.identity.v2
- Includes
 - PasswordMethod—uses username/password to authenticate
 - TokenMethod—uses existing token to authenticate
- auth_url must be keystone V2 root (http://url:5000/v2.0)

A V2 request can employ only one authentication method.

V3 identity plugins:

- defined in module keystoneclient.auth.identity.v3
- A V3 request can contain multiple authentication methods:
 - Includes "AuthMethod" objects that are sent to V3 keystone interface:
 - PasswordMethod—uses username/password to authenticate
 - TokenMethod—uses existing token to authenticate
- auth_url must be keystone V3 root (http://url:5000/v3)

The following method is used to pass an AuthMethod object to the Keystone auth plugin:

Note that you can also access this plugin similarly to how a V2 plugin is accessed, by doing specific calls that employ only a single authentication method. For example:

- keystoneclient.auth.identity.v3.Password
 auth_url, username, password,
 other arguments)
- keystoneclient.auth.identity.v3.Token(auth_url, token, other arguments)
- "other arguments" includes: user_domain_id, project_domain_id, project_id, project_name, etc.

For example:

```
v3.Password(
      auth url=...v3 keystone root url,
      username=...,
      password=...
       user domain name =...,
      project domain name=...,
      project id=...
      etc.
```

```
def authfunc():
      from keystoneauth1.identity import v3
      from keystoneauth1 import session
      from keystoneclient.v3 import client
      auth = v3.Password(user_domain_name='default',
username='admin',password='secret',project_domain_nam
e='default', project_name='admin',
auth_url='http://localhost:5000/v3')
      sess=session.Session(auth=auth)
      keystone=client.Client(session=sess)
      return keystone
```

In the dode from the previous page, this line imports the keystone python client:

from keystoneauth1.identity import v3

The following line imports the keystone session library:

from keystoneauth1 import session

and this line (further down in the code) creates a session using the given authorization (the authorization is stored in the "auth" object):

sess=session.Session(auth=auth)

This line imports the keystone client library for the v3 interface:

from keystoneclient.v3 import client

and this line creates a client that employs the previously created session:

keystone=client.Client(session=sess)

keystone client.v3.client does changes and queries by using managers. Some example managers:

- keystoneclient.v3.users.UserManager
 - Create, delete, get, list user(s), add/remove users to groups, check if user is in a group
- keystoneclient.v3.services.ServiceManager
 - Create, retrieve, list, update delete services on the server
- keystoneclient.v3.groups.GroupManager
 - Create, delete, get, list, update groups
- keystoneclient.v3.endpoints.EndpointManager
 - Create, retrieve, delete, list endpoints
- keystoneclient.v3.ec2.EC2Manager
 - Create access key/secret key pair
 - Delete access key/secret key pair
 - Retrieve access key/secret key pair for given access key
 - List access key/secret key pairs for a given user

openstack.session.Session manages:

- authenticator
- transport
- user profile

A session does the following:

- Inserts authentication tokens into requests
- Retrieves endpoints from the authenticator
- Determines service preferences
- Sends requests using the appropriate transport

Openstack.session.Session(profile, user-agent=..., various arguments)

So a session object is associated with a:

- profile
- User agent
 - A user agent string is used to identify a software agent (software acting on behalf of a user)
 - Could include operating system, application type, revision.

- A profile object contains user preferences for various services:
 - service name
 - Service type—compute, identity, object-store, etc.
 - Desired name of the service
 - Region associated with the service
 - Version of the service—could be V2, V3, etc.
 - Interface—interface associated with the specified service

The OpenStack python SDK works as follows:

- Connection methods (associated with the Connection interface) accept and return Resource objects (associated with the Resource interface)
- Your session, authentication, transport and profile are maintained by an instance of a Connection
 - Based on a particular user's profile, various OpenStack services are available to be accessed by that users through the Connection

Resources can be used directly (as well as being accessed through the Connection interface)

A Resource object is associated with the REST API of each service

The Resource class supports:

- Create Read Update Delete operations on the REST APIs
- Calling HTTP GET, POST, PUT on the associated Session
- Creating URLs

Value's sent to/received from the service are implemented as attributes (type openstack.resource.prop) on the Resource class:

 These attributes can include a type that can be validated when requests are received

The attribute base_path should correspond to the url of the resource

Other Resource attributes are checked before a request is made to the associated REST API:

- allow_create, allow_delete, allow_retrieve, allow_head, allow_list
- these each have values of true/false

OpenStack Python Interface (cont'd) A Proxy class is associated with each service

A Proxy class provides a high level interface for a user that:

- Manages a Session
- Works with lower level Resource objects

A Connection object provides a high level interface for a user that:

- Provides a Proxy object for each Resource object (associated with each service)
- Is "built on top" of a Session object